

[« Prev](#)[Next »](#)

Postgres Weekly

Issue 5

April 10, 2013

A big hello to new subscribers this week - thanks for joining us! As always, a big thanks to [Craig Kerstiens](#) for curating Postgres Weekly and if you want to recommend anything for a future issue, tweet [@craigkerstiens](#) and he might consider it! :-)

This week

[Top 10 Reasons I like Postgres Over SQL Server](#)

There's a variety of reasons we choose the databases we do. Rob Sullivan, a former Microsoft SQL Server DBA, here lays out 10 very clear reasons why he prefers Postgres over SQL Server.

[PostgreSQL: 2013-04-04 Security Release FAQ](#)

Postgres had a release on Thursday 4/4 to patch the worst security vulnerability for PostgreSQL since 2006. If you have not upgraded to the latest minor version do so immediately. The community has provided this clear FAQ around the specific vulnerability to answer many questions people may have.

[PostgreSQL Documentation \(devel\): JSON Functions and Operators](#)

Postgres 9.2 received the JSON datatype, but it was limited to validation and not much optimization for indexing and searching. JSON searching has now arrived in Postgres development builds and is likely to be available in Postgres 9.3. The docs show a preview of the functionality that is likely to be available.

[Calculating rolling cohort retention with SQL](#)

Calculating retention and understanding your data such as by doing cohort analysis is a valuable thing. It doesn't require a tool to always perform this, metrika shows how this can easily be done within SQL on the data you already have.

[Expression Indexes](#)

With any database of regular use developers will add indexes to them; with an application with a reasonable amount of data you may want your indexes to have a bit more intelligence to them. Here Bruce Momjian of Postgres Core explains how you can actually compute expressions within indexes – resulting in indexes that are intelligent as your application.

[PostgreSQL HyperLogLog](#)

PostgreSQL-hll or Postgres HyperLogLog which we previously featured has some great benefits. It allows for roughly 10 billion uniques to be counted within a 1280 byte datatype and have efficient queries across the dataset as well as support for [unions other operations](#). If dealing with high transaction data and you care about tracking such this is an extension to check out.

[PostgreSQL Participating in Google Summer of Code 2013](#)

Google will be funding several students to work with mentors from the PostgreSQL project in order to hack on the project over the summer.

Did you know...

[We publish several other newsletters like Postgres Weekly](#)

Are you interested in Ruby, JavaScript, HTML5, browser technology, Dart, or just programming in general? We publish a variety of other newsletters like [HTML5 Weekly](#), [Ruby Weekly](#) and [The Modern Web Observer](#) too :-)

[« Prev](#)[Next »](#)[Subscribe now »](#)

Easy to unsubscribe at any time. Your e-mail address [is safe](#) — here's our [privacy policy](#).