

[« Prev](#)[Next »](#)

Postgres Weekly

Issue 29

September 25, 2013

Featured

[Postgres Open 2013 Wrapup](#)

Postgres Open was last week in Chicago, with talks targeting a broad range of users, many delivered from major contributors to Postgres. While videos are still to come, many of the slides from the conference are becoming available.

[POSTGRESQL WKI](#)

[PostgreSQL 9.3 and JSON \[PDF\]](#)

JSON support in PostgreSQL 9.3 is much improved from PostgreSQL 9.2. One of the major contributors to the JSON support was at Postgres Open in Chicago last week and gave a great deeper dive into what changes came along with 9.3.

[ANDREW DUNSTAN](#)

From our sponsor

[Heroku Postgres - Fully managed Postgres on EC2](#)

There's a lot to worry about when building your application; ensuring your database is secured, backed up, available shouldn't have to be part of that. Let Heroku Postgres focus on keeping your database running so you can take advantage of all the awesome features within Postgres.

Heroku Postgres is the worlds largest provider of Postgres as a Service, offering fully managed databases on top of EC2. Get started today for free and scale easily as your application needs to.

The rest

[Returning Hierarchical Data in a Single SQL Query](#)

Hierarchical data is common in a lot of applications, working with and reporting against it isn't always as easy as getting it there in the first place. Taking advantage of CTEs (Common Table Expressions) which are still a heavily underrated feature can make it much easier.

[ANDREW BENDER](#)

[Using JSON data type in production with Socorro](#)

JSON data is everywhere, but those using it in large applications within Postgres may still be a pretty small number as support only came in Postgres 9.2. With it continuing to improve in 9.3 we'll see more adopt, and having large applications such as Socorro by Mozilla take advantage of it will on further validation of its value.

[SELENA DECKELMANN](#)

[Federating Queries Using postgres_fdw \[PDF\]](#)

Foreign data wrappers promise to be a great improvement now and going forward for Postgres. Whether its dealing with disparate data sources or looking to build a full federated database they make it much easier. Here's one of two talks at Postgres Open which dig into taking advantage of the Postgres FDW available in 9.3.

[Scalable Uniques in Postgres with HLL](#)

If you've ever done a `SELECT count()` in Postgres across a large dataset you've experienced the pain of how slow this can be.

HyperLogLog is an algorithm for cardinality estimation which can get very close to `count()` with a much smaller footprint. Here's slides from a brief talk by myself at Postgres Open on the PostgreSQL extension for HLL.

[« Prev](#)[Next »](#)

Subscribe now »

Easy to unsubscribe at any time. Your e-mail address [is safe](#) — [here's our privacy policy](#).