

[« Prev](#)[Next »](#)

Postgres Weekly

Issue 71

September 3, 2014

[PGConf.EU 2014 schedule posted: Matt Asay to deliver keynote](#)

PGConf EU is happening in Madrid this year which sounds like a great location, in addition there's a [great lineup](#) of talks. Including what makes to be an interesting keynote from Matt Asay, Vice President of Community at MongoDB.

[POSTGRESQL.ORG](#)

[Postgres multicolumn indexes to save the day](#)

Indexes are great things that will make your database perform better, but in Postgres you have a few options from functional and conditional ones, to the different types - BTree, GIN, GiST, KNN, SP-GiST. Multicolumn indexes while maybe not your most commonly used, when needed can be extremely effective.

[JULIEN VEHENT](#)

[Postgres Skills are in Demand - Add to your skill set and get certified](#)

1000+ jobs are posted each month listing Postgres skills – enhance your talents and get certified with expert training courses. The EDB certification program sets the global standard for Postgres professionals, and is definitive proof of your Postgres skills.

[ENTERPRISEDB](#) **SPONSORED**

[Replicated PostgreSQL with pgpool2](#)

If you're running a Postgres installation yourself you want some form of replication setup. Of course there's lots of options such as [WAL-E](#) and [Barman](#), and then there's PG Pool which handles quite a few things beyond just replication. Here's your perfect guide for getting at least the replication and failover setup.

[MICHAEL STAPELBERG](#)

[Waiting for Postgres 9.5: Add min and max aggregates for inet/cidr data types](#)

Using inet within Postgres? If you've got the specific use case of wanting to find the min and max in some grouping, it's going to be available in 9.5. While likely not earth shattering with each new function, Postgres becomes even more user friendly.

[HUBERT DEPESZ LUBACZEWSKI](#)

[Streaming directly into Postgres with Hapi.js and pg-copy-stream](#)

The copy command is extremely handy for bulk loading data, if you're doing anything with inserting lots of data it's what you want to use. Here's an awesome example of a Node app taking advantage of just that and getting a 10x performance improvement over plain inserts.

[GRAHAM JENSON](#)

[Many to many relationships with Postgresql's hstore in Rails 4](#)

Hstore is awesome and powerful in Postgres, and the more people take advantage of it in their [frameworks](#) like [Django](#) and [Rails](#) the more benefit you find. Here's the latest example going even further providing many to many relationships directly in rails with ActiveRecord.

[PIER-OLIVIER THIBAUT](#)[« Prev](#)[Next »](#)[Subscribe now »](#)

Easy to unsubscribe at any time. Your e-mail address [is safe](#) — [here's our privacy policy](#).